Informed search algorithms

This lecture topic Chapter 3.3

Outline

- Review limitations of uninformed search methods
- Informed (or heuristic) search uses problem-specific heuristics to improve efficiency
 - Best-first, A* (and if needed for memory limits, RBFS, SMA*)
 - Techniques for generating heuristics
 - A* is optimal with admissible (tree)/consistent (graph) heuristics
 - A* is quick and easy to code, and often works *very* well

Heuristics

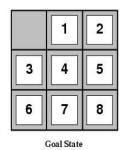
- A structured way to add "smarts" to your solution
- Provide *significant* speed-ups in practice
- Still have worst-case exponential time complexity
- In AI, "NP-Complete" means "Formally interesting"

Limitations of uninformed search

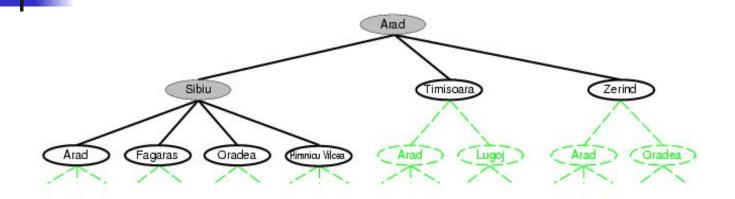
- Search Space Size makes search tedious
 - Combinatorial Explosion
- For example, 8-puzzle
 - Avg. solution cost is about 22 steps
 - branching factor ~ 3
 - Exhaustive search to depth 22:
 - 3.1 x 10¹⁰ states
 - E.g., d=12, IDS expands 3.6 million states on average

[24 puzzle has 10^{24} states (much worse)]

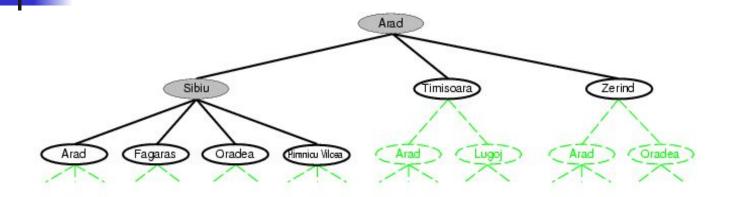




Recall tree search...



Recall tree search...



function TREE-SEARCH(problem, strategy) returns a solut initialize the search tree using the initial state of problem loop do

This "strategy" is what differentiates different search algorithms

if there are no candidates for expansion then return failure choose a leaf node for expansion according to *strategy*if the node contains a goal state then return the corresponding solution else expand the node and add the resulting nodes to the search tree

Heuristic search

- Idea: use an evaluation function f(n) for each node and a heuristic function h(n) for each node
 - g(n) = known path cost so far to node n.
 - $\ddot{h}(n) = \underline{\text{estimate}}$ of (optimal) cost to goal from node n.
 - f(n) = g(n) + h(n) = estimate of total cost to goal through node n.
 - f(n) provides an <u>estimate</u> for the total cost:
 - Expand the node n with smallest f(n).
- Implementation:
 Order the nodes in frontier by increasing estimated cost.
- Evaluation function is an <u>estimate</u> of node quality
 - More accurate name for "best first" search would be "seemingly best-first search"
- Search efficiency depends on heuristic quality!
 - → The better your heuristic, the faster your search!

Heuristic function

Heuristic:

- Definition: a commonsense rule (or set of rules) intended to increase the probability of solving some problem
- Same linguistic root as "Eureka" = "I have found it"
- "using rules of thumb to find answers"

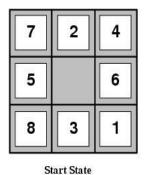
Heuristic function h(n)

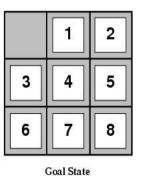
- Estimate of (optimal) remaining cost from n to goal
- Defined using only the <u>state</u> of node n
- h(n) = 0 if n is a goal node
- Example: straight line distance from n to Bucharest
 - Note that this is not the true state-space distance
 - It is an estimate actual state-space distance can be higher
- Provides problem-specific knowledge to the search algorithm

Heuristic functions for 8-puzzle



- Avg. solution cost is about 22 steps
- branching factor ~ 3
- Exhaustive search to depth 22:
 - 3.1 x 10¹⁰ states.

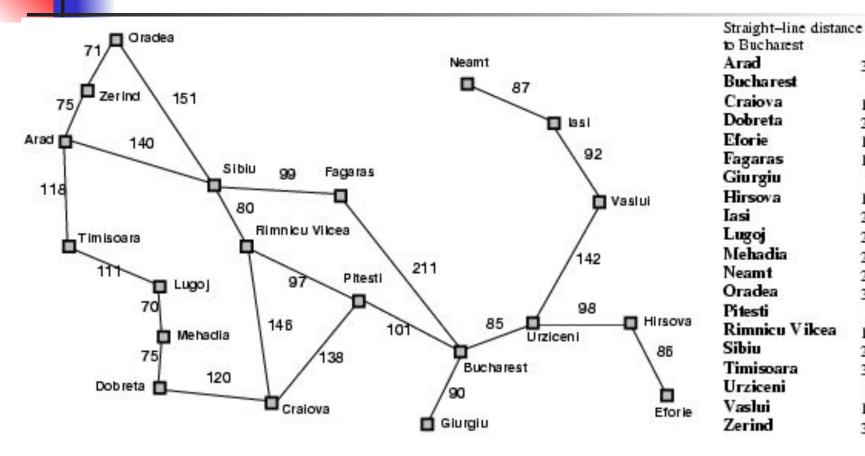




- A good heuristic function can reduce the search process.
- Two commonly used heuristics
 - h_1 = the number of misplaced tiles
 - $h_1(s)=8$
 - h_2 = the sum of the distances of the tiles from their goal positions (Manhattan distance).

•
$$h_2(s)=3+1+2+2+3+3+2=18$$

Romania with straight-line dist.



Relationship of Search Algorithms

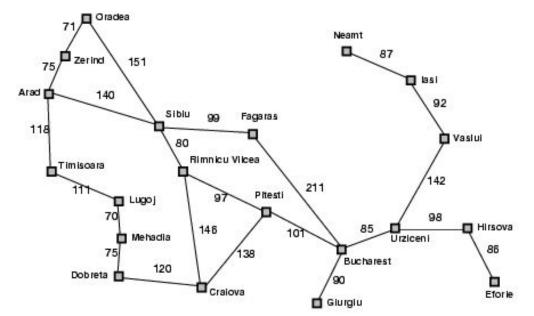
- g(n) = known cost so far to reach n
- h(n) =estimated (optimal) cost from n to goal
- f(n) = g(n) + h(n)
 - = estimated (optimal) total cost of path through *n* to goal
- Uniform Cost search sorts frontier by g(n)
- Greedy Best First search sorts frontier by h(n)
- A* search sorts frontier by f(n)
 - Optimal for admissible/consistent heuristics
 - Generally the preferred heuristic search
- Memory-efficient versions of A* are available
 - RBFS, SMA*

Greedy best-first search (often called just "best-first")

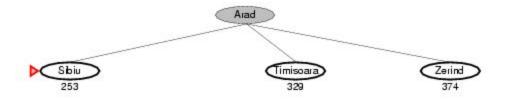
- h(n) = estimate of cost from n to goal
 - e.g., h(n) = straight-line distance from n to Bucharest

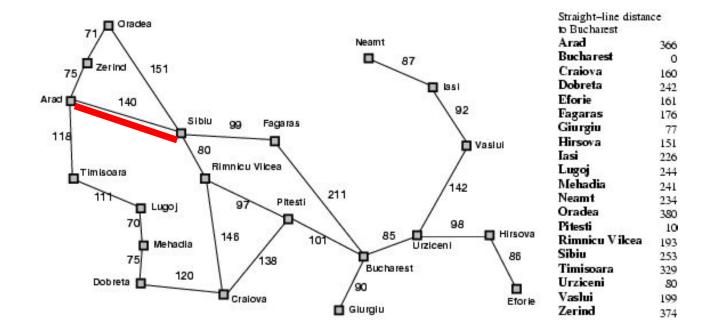
- Greedy best-first search expands the node that appears to be closest to goal.
 - Priority queue sort function = h(n)








Straight-line distan	ce
to Bucharest	
Arad	36
Bucharest	(
Craiova	160
Dobreta	24
Eforie	16
Fagaras	176
Giurgiu	7
Hirsova	15
Iasi	220
Lugoj	24
Mehadia	24
Neamt	23
Oradea	39
Pitesti	10
Rimnicu Vilcea	19
Sibiu	25
Timisoara	32
Urziceni	8
Vaslui	19
Zerind	37












Optimal Path



Properties of greedy best-first search

- Complete?
 - Tree version can get stuck in loops.
 - Graph version is complete in finite spaces.
- <u>Time?</u> *O*(*b*^{*m*})
 - A good heuristic can give dramatic improvement
- Space? O(b^m)
 - Keeps all nodes in memory
- Optimal? No

e.g., Arad □ Sibiu □ Rimnicu Vilcea □ Pitesti □ Bucharest is shorter!

A* search

- Idea: avoid paths that are already expensive
 - Generally the preferred simple heuristic search
 - Optimal if heuristic is: admissible(tree)/consistent(graph)
- Evaluation function f(n) = g(n) + h(n)
 - g(n) = known path cost so far to node n.
 - h(n) = estimate of (optimal) cost to goal from node n.
 - f(n) = g(n)+h(n)= <u>estimate</u> of total cost to goal through node n.
- Priority queue sort function = f(n)

Admissible heuristics

- A heuristic h(n) is admissible if for every node n,
 h(n) ≤ h*(n),
 where h*(n) is the true cost to reach the goal state from n.
- An admissible heuristic never overestimates the cost to reach the goal, i.e., it is optimistic (or at least, never pessimistic)
 - Example: $h_{SID}(n)$ (never overestimates actual road distance)
- Theorem:

If h(n) is admissible, A^* using TREE-SEARCH is optimal

Admissible heuristics

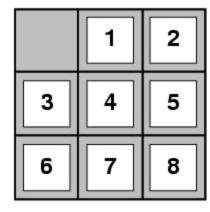
E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1





Goal State

•
$$h_1(S) = ?$$

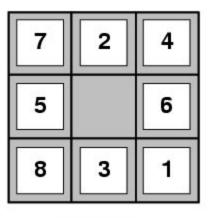
•
$$h_2(S) = ?$$

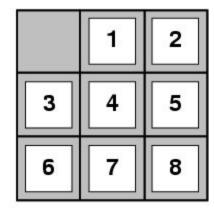
Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)





•
$$h_1(S) = ?8$$

Start State

Goal State

•
$$h_2(S) = ? 3+1+2+2+3+3+2 = 18$$

Consistent heuristics (consistent => admissible)

A heuristic is consistent if for every node n, every successor n' of n generated by any action a,

$$h(n) \le c(n,a,n') + h(n')$$

If h is consistent, we have

$$f(n') = g(n') + h(n')$$
 (by def.)
= $g(n) + c(n,a,n') + h(n')$ ($g(n')=g(n)+c(n.a.n')$)
 $\geq g(n) + h(n) = f(n)$ (consistency)
 $f(n') \geq f(n)$

i.e., f(n) is non-decreasing along any path.

It's the triangle inequality!

**Reeps all checked nodes in memory to avoid repeated If h(n) is consistent, A* using GRAPH-SEARCH is optimal states

Admissible (Tree Search) vs. Consistent (Graph Search)

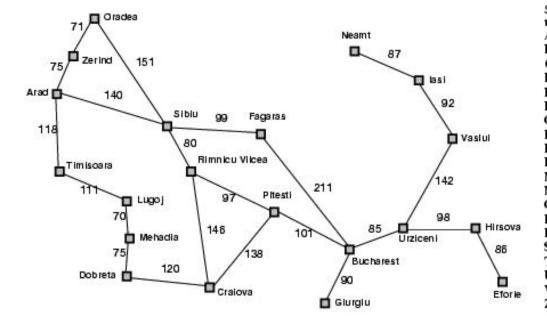


- In graph search you often find a long cheap path to a node after a short expensive one, so you might have to update all of its descendants to use the new cheaper path cost so far
- A consistent heuristic avoids this problem (it can't happen)
- Consistent is slightly stronger than admissible
- Almost all admissible heuristics are also consistent

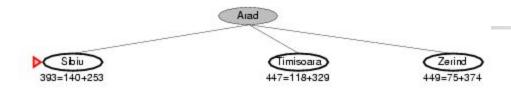
Could we do optimal graph search with an admissible heuristic?

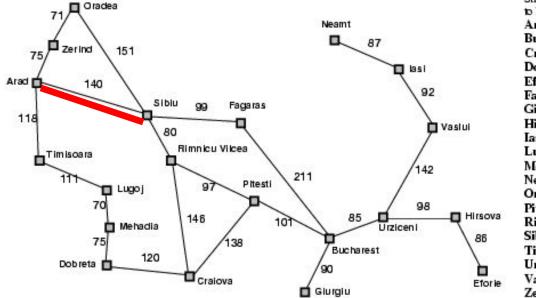
- Yes, but you would have to do additional work to update descendants when a cheaper path to a node is found
- A consistent heuristic avoids this problem



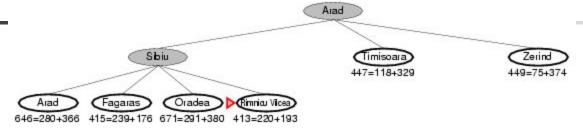


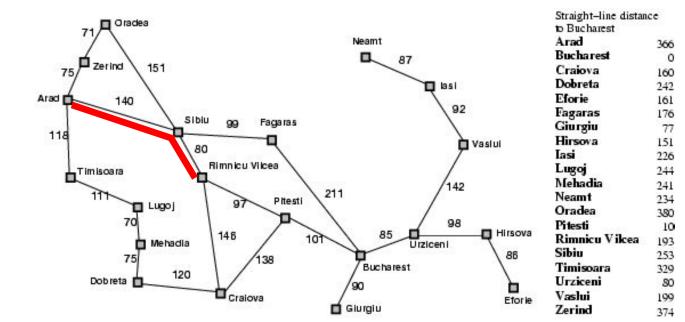
ce
360
(
160
242
16
170
7
15
226
24
24
23-
380
10
193
25
329
80
199
37-








Straight-line distance to Bucharest Arad 366 Bucharest 0 Craiova 160 Dobreta 242 Eforie 161 Fagaras 176 Giurgiu 77 Hirsova 151 Iasi 226 Lugoj 244 Mehadia 241 Neamt 234 Oradea 380 Pitesti 10 Rimnicu Vilcea 193 Sibiu 253 Timisoara 329 Urziceni 80 Vaslui 199 Zerind 374

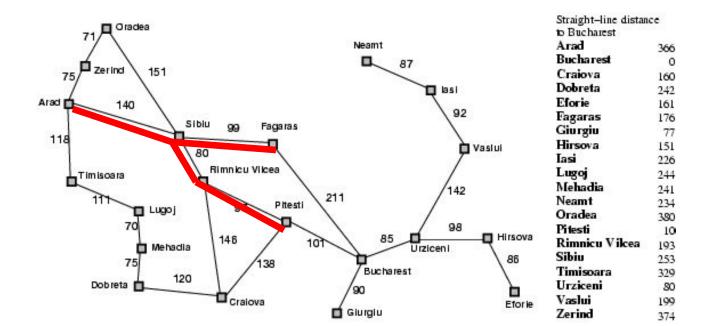


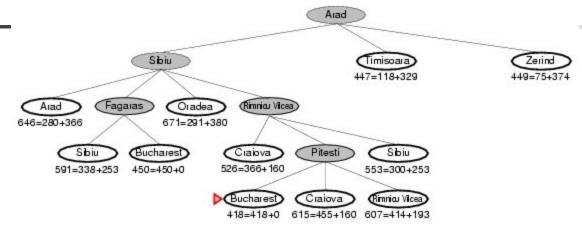


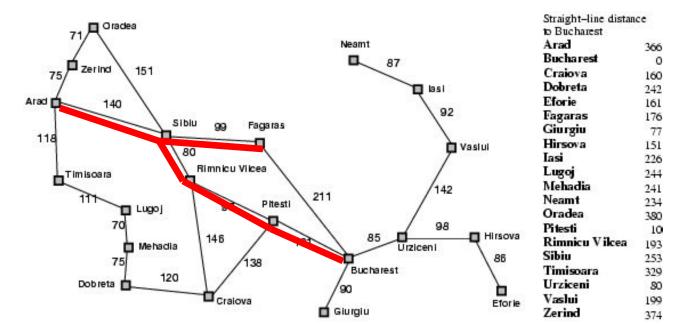













Contours of A* Search

- A* expands nodes in order of increasing f value
- Gradually adds "f-contours" of nodes
- Contour *i* has all nodes with $f=f_{j}$ where $f_{i} < f_{i+1}$




Properties of A*

- Complete? Yes
 - (unless there are infinitely many nodes with $f \le f(G)$; can't happen if step-cost $\ge \varepsilon > 0$)
- <u>Time/Space?</u> Exponential O(b^d)
 - except if: $|h(n)-h^*(n)| \leq O(\log h^*(n))$
- Optimal? Yes
 - (with: Tree-Search, admissible heuristic; Graph-Search, consistent heuristic)
- Optimally Efficient? Yes
 - (no optimal algorithm with same heuristic is guaranteed to expand fewer nodes)

Memory Bounded Heuristic Search: Recursive Best First Search (RBFS)

• How can we solve the memory problem for A* search?


- Idea: Try something like depth first search, but let's not forget everything about the branches we have partially explored.
- We remember the best f(n) value we have found so far in the branch we are deleting.



RBFS changes its mind very often in practice.

This is because the f=g+h become more accurate (less optimistic) as we approach the goal. Hence, higher level nodes have smaller f-values and will be explored first.

Problem: We should keep in memory whatever we can.




Simple Memory Bounded A* (SMA*)

- This is like A*, but when memory is full we delete the worst node (largest f-value).
- Like RBFS, we remember the best descendent in the branch we delete.
- If there is a tie (equal f-values) we delete the oldest nodes first.
- simple-MBA* finds the optimal *reachable* solution given the memory constraint.

 A Solution is not reachable
- Time can still be exponential.

A Solution is not reachable if a single path from root to goal does not fit into memory

SMA* pseudocode (not in 2nd edition of R&N)



```
n \square deepest least-f-cost node in Queue
   if GOAL-TEST(n) then return success
   s \square NEXT-SUCCESSOR(n)
   if s is not a goal and is at maximum depth then
      f(s) \square \infty
   else
      f(s) \square MAX(f(n),g(s)+h(s))
   if all of n's successors have been generated then
      update n's f-cost and those of its ancestors if necessary
   if SUCCESSORS(n) all in memory then remove n from Queue
   if memory is full then
      delete shallowest, highest-f-cost node in Queue
      remove it from its parent's successor list
      insert its parent on Queue if necessary
   insert s in Queue
end
```

Simple Memory-bounded A* (SMA*)

(Example with 3-node memory)

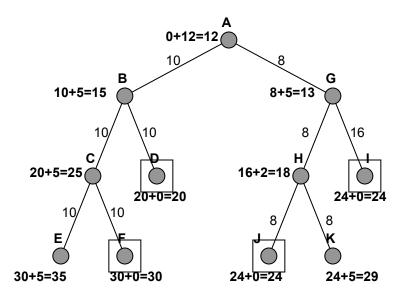
maximal depth is 3, since memory limit is 3. This branch is now useless.

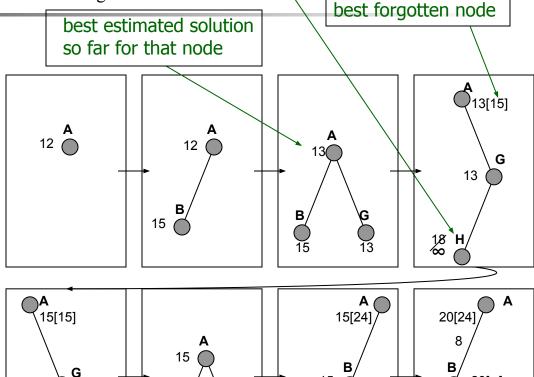
20[∞]

Progress of SMA*. Each node is labeled with its *current f*-cost. Values in parentheses show the value of the best forgotten descendant.

Search space

$$g+h=f$$
 $\square = goal$





Algorithm can tell you when best solution found within memory constraint is optimal or not.

24

24[∞]

Memory Bounded A* Search

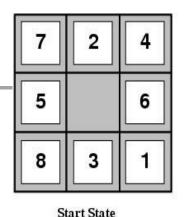
- The Memory Bounded A* Search is the best of the search algorithms we have seen so far. It uses all its memory to avoid double work and uses smart heuristics to first descend into promising branches of the search-tree.
- If memory not a problem, then plain A* search is easy to code and performs well.

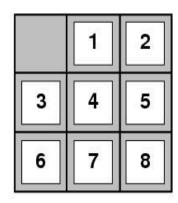
Heuristic functions



- 8-puzzle
 - Avg. solution cost is about 22 steps
 - branching factor ~ 3
 - Exhaustive search to depth 22:
 - 3.1 x 10¹⁰ states.
 - A good heuristic function can reduce the search process.
- Two commonly used heuristics
 - h_1 = the number of misplaced tiles
 - $h_1(s)=8$
 - h_2 = the sum of the axis-parallel distances of the tiles from their goal positions (manhattan distance).

•
$$h_2(s)=3+1+2+2+3+3+2=18$$





Goal State

Dominance

- IF $h_2(n) \ge h_1(n)$ for all n (both admissible) THEN h_2 dominates h_1
 - h_2 is always better for search than h_1
 - h_2 guarantees to expand no more nodes than does h_1
 - h_2 almost always expands fewer nodes than does h_1
- Typical 8-puzzle search costs (average number of nodes expanded):
 - d=12IDS = 3,644,035 nodes $A^*(h_1) = 227$ nodes $A^*(h_2) = 73$ nodes
 - d=24 ÎDS = too many nodes $A^*(h_1) = 39,135$ nodes $A^*(h_2) = 1,641$ nodes

Effective branching factor: b*

- Let A* generate N nodes to find a goal at depth d
 - b^* is the branching factor that a uniform tree of depth d would have in order to contain N+1 nodes.

$$N+1 = 1+b^* + (b^*)^2 + ... + (b^*)^d$$

$$N+1 = ((b^*)^{d+1} - 1)/(b^* - 1)$$

$$N \approx (b^*)^d \implies b^* \approx \sqrt[d]{N}$$

- For sufficiently hard problems, the measure b* usually is fairly constant across different problem instances.
- A good guide to the heuristic's overall usefulness.
- A good way to compare different heuristics.

Effective Branching Factor Pseudo-code (Binary search)

```
PROCEDURE EFFBRANCH (START, END, N, D, DELTA)
 COMMENT DELTA IS A SMALL POSITIVE NUMBER FOR ACCURACY OF RESULT.
 MID := (START + END) / 2.
 IF (END - START < DELTA)
   THEN RETURN (MID).
 TEST := EFFPOLY (MID, D).
 IF (TEST < N+1)
   THEN RETURN (EFFBRANCH (MID, END, N, D, DELTA) )
   ELSE RETURN (EFFBRANCH (START, MID, N, D, DELTA) ).
END EFFBRANCH.
PROCEDURE EFFPOLY (B, D)
 ANSWER = 1.
 TEMP = 1.
 FOR I FROM 1 TO (D-1) DO
   TEMP := TEMP * B.
   ANSWER := ANSWER + TEMP.
 ENDDO.
 RETURN (ANSWER).
END EFFPOLY.
```

- For binary search please see: http://en.wikipedia.org/wiki/Binary search algorithm
- An attractive alternative is to use Newton's Method (next lecture) to solve for the root (i.e., f(b)=0) of $f(b) = 1 + b + ... + b^d (N+1)$

Effectiveness of different heuristics

d	Search Cost		Effective Branching Factor			
	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	2 000 n-000	539	113	The Tell	1.44	1.23
16	eserve same	1301	211	10,000	1.45	1.25
18	_	3056	363		1.46	1.26
20		7276	676		1.47	1.27
22		18094	1219	Daxed = satt	1.48	1.28
24	-	39135	1641	a nes a eny v	1.48	1.26

 Results averaged over random instances of the 8-puzzle

Inventing heuristics via "relaxed problems"

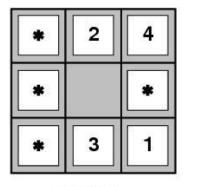
- A problem with fewer restrictions on the actions is called a relaxed problem
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then $h_1(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to any adjacent square, then $h_2(n)$ gives the shortest solution
- Can be a useful way to generate heuristics
 - E.g., ABSOLVER (Prieditis, 1993) discovered the first useful heuristic for the Rubik's cube puzzle

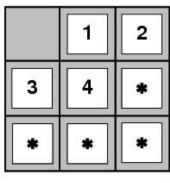
More on heuristics

- $h(n) = max\{ h_1(n), h_2(n), ..., h_k(n) \}$
 - Assume all h functions are admissible
 - E.g., h1(n) = # of misplaced tiles
 - E.g., h2(n) = manhattan distance, etc.
 - max chooses least optimistic heuristic (most accurate) at each node
- $h(n) = w_1 h_1(n) + w_2 h_2(n) + ... + w_k h_k(n)$
 - A convex combination of features
 - Weighted sum of h(n)'s, where weights sum to 1
 - Weights learned via repeated puzzle-solving
 - Try to identify which features are predictive of path cost

Pattern databases

- Admissible heuristics can also be derived from the solution cost of a subproblem of a given problem.
- This cost is a lower bound on the cost of the real problem.
- Pattern databases store the exact solution to for every possible subproblem instance.
 - The complete heuristic is constructed using the patterns in the DB





Start State

Goal State

Summary

- Uninformed search methods have uses, also severe limitations
- Heuristics are a structured way to add "smarts" to your search
- Informed (or heuristic) search uses problem-specific heuristics to improve efficiency
 - Best-first, A* (and if needed for memory limits, RBFS, SMA*)
 - Techniques for generating heuristics
 - A* is optimal with admissible (tree)/consistent (graph) heuristics
- Can provide significant speed-ups in practice
 - E.g., on 8-puzzle, speed-up is dramatic
 - Still have worst-case exponential time complexity
 - In AI, "NP-Complete" means "Formally interesting"
- Next lecture topic: local search techniques
 - Hill-climbing, genetic algorithms, simulated annealing, etc.